

A Binary Search Tree Implementation

Chapter 25

Data Structures and Abstractions with Java, 4e, Global Edition
Frank Carrano

Binary Search Tree (BST)

- For each node in a binary search tree
 - Node's data is greater than all data in node's left subtree
 - Node's data is less than all data in node's right subtree
- Every node in a binary search tree is the root of a binary search tree

Binary Search Tree




FIGURE 25-1 A binary search tree of names

Interface for the BST

```
package TreePackage;
import java.util.Iterator;
public interface SearchTreeInterface<T extends Comparable<? super T>>
{
    /** Searches for a specific entry in this tree.
     * @param entry An object to be found.
     * @return True if the object was found in the tree. */
    public boolean contains(T entry);

    /** Retrieves a specific entry in this tree.
     * @param entry An object to be found.
     * @return Either the object that was found in the tree or
     *         null if no such object exists. */
    public T getEntry(T entry);

    /** Adds a new entry to this tree, if it does not match an existing
     *    object in the tree. Otherwise, replaces the existing object with
     *    the new entry.
```

Interface for the BST

```
/** Adds a new entry to this tree, if it does not match an existing
    object in the tree. Otherwise, replaces the existing object with
    the new entry.
    @param newEntry An object to be added to the tree.
    @return Either null if newEntry was not in the tree already, or
            the existing entry that matched the parameter newEntry
            and has been replaced in the tree. */
public T add(T newEntry);

/** Removes a specific entry from this tree.
    @param entry An object to be removed.
    @return Either the object that was removed from the tree or
            null if no such object exists. */
public T remove(T entry);

/** Creates an iterator that traverses all entries in this tree.
    @return An iterator that provides sequential and ordered access
            to the entries in the tree. */
public Iterator<T> getInorderIterator();
```

Understanding the Specifications

- Methods will use return values instead of exceptions to indicate whether an operation has failed.
- Pay attention to
 - the entry's `equals` method
 - the entry's `compareTo` method

Understanding the Specifications

(a) Before `myTree.add(whitney2)` executes




FIGURE 25-2 Adding an entry that matches an entry already in a binary search tree

Understanding the Specifications

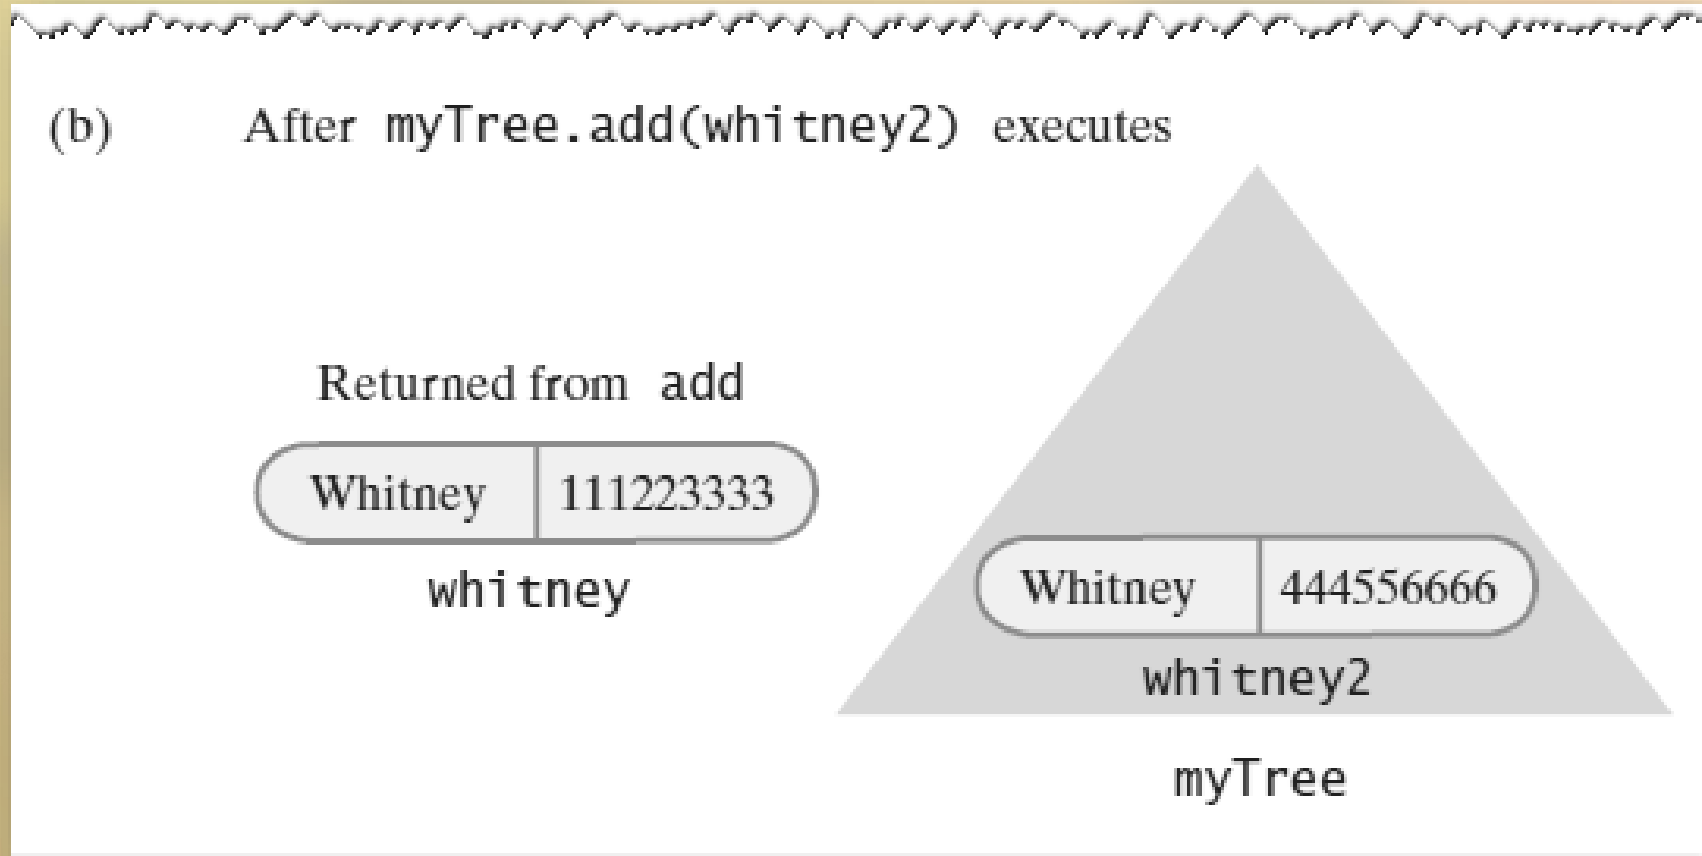


FIGURE 25-2 Adding an entry that matches an entry already in a binary search tree

If Duplicate Entries Are Allowed

- If any entry e has a duplicate entry d , we require that d occur in the right subtree of e 's node
- For each node in a binary search tree:
 - Data in a node is greater than data in node's left subtree
 - Data in a node is less than *or equal to* data in node's right subtree

If Duplicate Entries Are Allowed




FIGURE 25-3 A binary search tree with duplicate entries

Searching and Retrieving

Algorithm **bstSearch(binarySearchTreeRoot, desiredObject)**

// Searches a binary search tree for a given object.

// Returns true if the object is found.

if (binarySearchTreeRoot *is null*)

return false

else if (desiredObject == *object in* binarySearchTreeRoot)

return true

else if (desiredObject < *object in* binarySearchTreeRoot)

return bstSearch(*left child of* binarySearchTreeRoot, desiredObject)

else

return bstSearch(*right child of* binarySearchTreeRoot, desiredObject)

Algorithm that describes actual implementation more closely

```

public T getEntry(T entry)
{
    return findEntry(getRootNode(), entry);
} // end getEntry

private T findEntry(BinaryNode<T> rootNode, T entry)
{
    T result = null;

    if (rootNode != null)
    {
        T rootEntry = rootNode.getData();

        if (entry.equals(rootEntry))
            result = rootEntry;
        else if (entry.compareTo(rootEntry) < 0)
            result = findEntry(rootNode.getLeftChild(), entry);
        else
            result = findEntry(rootNode.getRightChild(), entry);
    } // end if

    return result;
} // end findEntry

```

The method **getEntry** uses **findEntry**

Searching and Retrieving

```
public boolean contains(T entry)
{
    return getEntry(entry) != null;
} // end contains
```

Method **contains** can simply call **getEntry** to see whether a given entry is in the tree:

Adding an Entry




FIGURE 25-4 (a) A binary search tree;

Adding an Entry

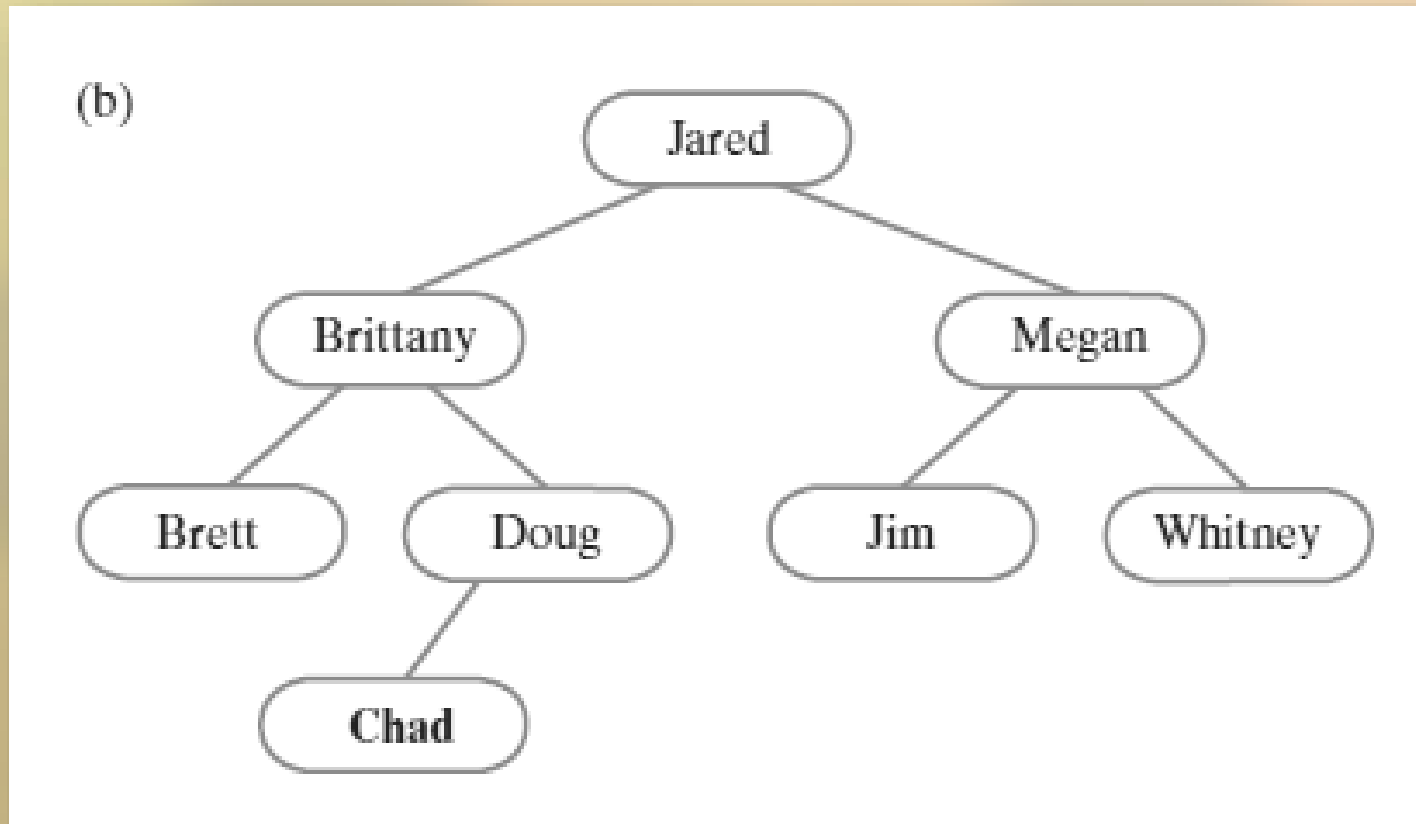


FIGURE 25-4 (b) the same tree after adding *Chad*

Recursive Implementation




FIGURE 25-5 Recursively adding *Chad* to smaller subtrees of a binary search tree

Recursive Implementation




FIGURE 25-5 Recursively adding *Chad* to smaller subtrees of a binary search tree

Recursive Implementation

Algorithm **add(binarySearchTree, newEntry)**

// Adds a new entry to a binary search tree.

*// Returns null if newEntry did not exist already in the tree. Otherwise, returns the
// tree entry that matched and was replaced by newEntry.*

result = null

if (binarySearchTree is empty)

Create a node containing newEntry and make it the root of binarySearchTree

else

result = addEntry(binarySearchTree, newEntry)

return result;

Handle the addition to an empty binary
search tree as a special case

Recursive Implementation

Algorithm addEntry(binarySearchTree, newEntry)

// Adds a new entry to a binary search tree that is not empty.

// Returns null if newEntry did not exist already in the tree. Otherwise, returns the

// tree entry that matched and was replaced by newEntry.

result = null

if (newEntry matches the entry in the root of binarySearchTree)

{

 result = entry in the root

 Replace entry in the root with newEntry

}

else if (newEntry < entry in the root of binarySearchTree)

{

 if (the root of binarySearchTree has a left child)

 result = addEntry(left subtree of binarySearchTree, newEntry)

 else

 Give the root a left child containing newEntry

}

else // newEntry > entry in the root of binarySearchTree

Recursive algorithm for adding a new entry

Recursive Implementation

```
else // newEntry > entry in the root of binarySearchTree
{
    if (the root of binarySearchTree has a right child)
        result = addEntry(right subtree of binarySearchTree, newEntry)
    else
        Give the root a right child containing newEntry
}

return result
```

Recursive algorithm for adding a new entry

Recursive Implementation

```
public T add(T newEntry)
{
    T result = null;

    if (isEmpty())
        setRootNode(new BinaryNode<>(newEntry));
    else
        result = addEntry(getRootNode(), newEntry);

    return result;
} // end add
```

Recursive Implementation

```
// Adds newEntry to the nonempty subtree rooted at rootNode.  
private T addEntry(BinaryNode<T> rootNode, T newEntry)  
{  
    assert rootNode != null;  
    T result = null;  
    int comparison = newEntry.compareTo(rootNode.getData());  
  
    if (comparison == 0)  
    {  
        result = rootNode.getData();  
        rootNode.setData(newEntry);  
    }  
    else if (comparison < 0)  
    {  
        if (rootNode.hasLeftChild())  
            result = addEntry(rootNode.getLeftChild(), newEntry);  
        else  
            rootNode.setLeftChild(new BinaryNode<>(newEntry));  
    }  
}
```

Recursive Implementation

```
else
{
    assert comparison > 0;

    if (rootNode.hasRightChild())
        result = addEntry(rootNode.getRightChild(), newEntry);
    else
        rootNode.setRightChild(new BinaryNode<>(newEntry));
} // end if

return result;
} // end addEntry
```

Iterative Implementation

Algorithm addEntry(binarySearchTree, newEntry)

// Adds a new entry to a binary search tree that is not empty.

// Returns null if newEntry did not exist already in the tree. Otherwise, returns

// tree entry that matched and was replaced by newEntry.

result = null

currentNode = *root node of* binarySearchTree

found = false

while (found is false)

{

 if (newEntry matches the entry in currentNode)

 {

 found = true

 result = *entry in* currentNode

Replace entry in currentNode with newEntry

 }

Iterative Implementation

```
}  
else if (newEntry < entry in currentNode)  
{  
    if (currentNode has a left child)  
        currentNode = left child of currentNode  
    else  
    {  
        found = true  
        Give currentNode a left child containing newEntry  
    }  
}  
else // newEntry > entry in currentNode  
{
```

Iterative Implementation

```
else // newEntry > entry in currentNode
{
    if (currentNode has a right child)
        currentNode = right child of currentNode
    else
    {
        found = true
        Give currentNode a right child containing newEntry
    }
}
}

return result
```

Iterative Implementation

```
private T addEntry(T newEntry)
{
    BinaryNodeInterface<T> currentNode = getRootNode();
    assert currentNode != null;
    T result = null;
    boolean found = false;

    while (!found)
    {
        T currentEntry = currentNode.getData();
        int comparison = newEntry.compareTo(currentEntry);

        if (comparison == 0)
        {
            // newEntry matches currentEntry;
            // return and replace currentEntry
            found = true;
            result = currentEntry;
            currentNode.setData(newEntry);
        }
    }
}
```

Iterative Implementation

```
        currentNode.setLeftChild(new BinaryNode<T>(newEntry));
    } // end if
}
else
{
    assert comparison > 0;

    if (currentNode.hasRightChild())
        currentNode = currentNode.getRightChild();
    else
    {
        found = true;
        currentNode.setRightChild(new BinaryNode<T>(newEntry));
    } // end if
} // end if
} // end while

return result;
} // end addEntry
```

Removing an Entry Whose Node Is a Leaf

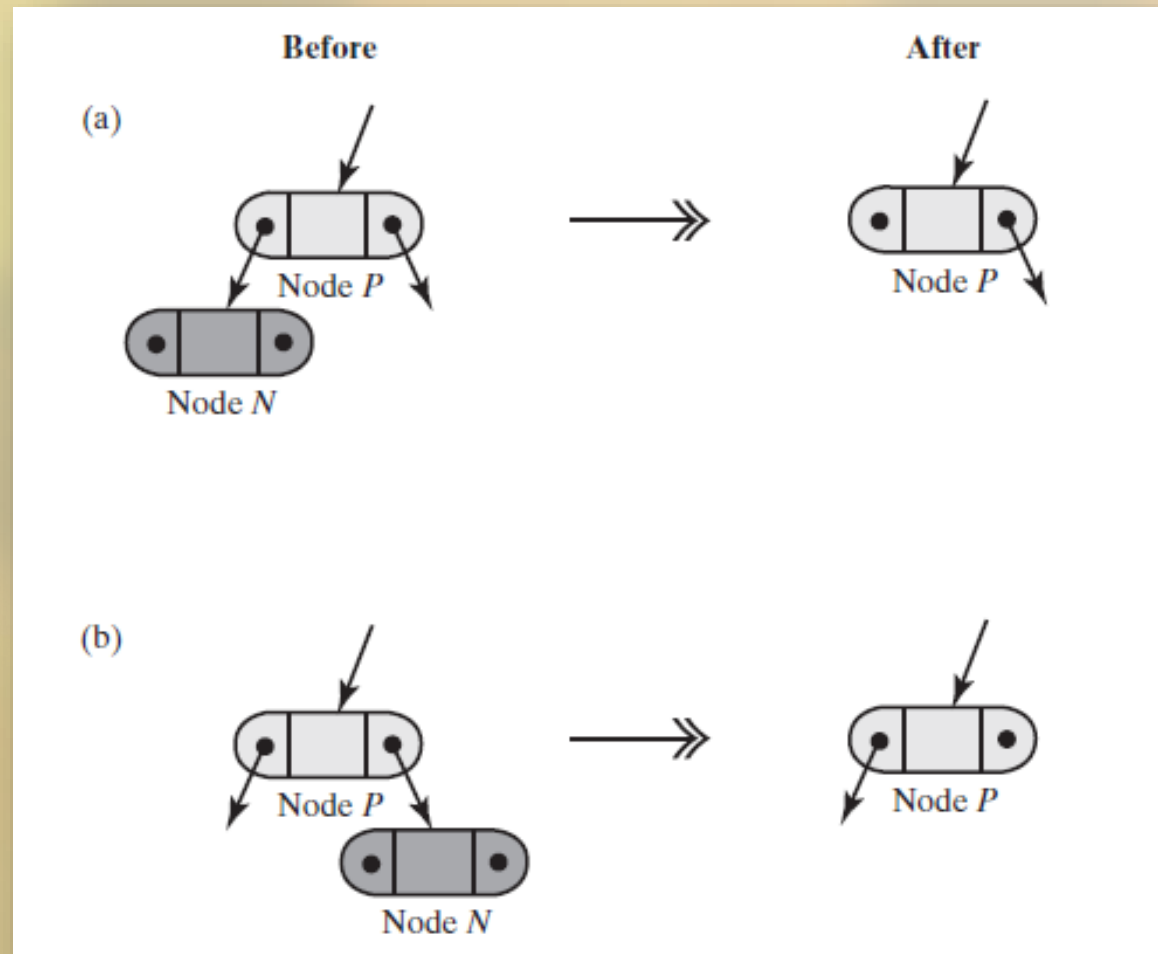


FIGURE 25-6 (a) Two possible configurations of a leaf node *N*; (b) the resulting two possible configurations after removing node

Removing an Entry Whose Node Has One Child

Two possible configurations before removal

After removal




FIGURE 25-7 Removing a node N from its parent node P when N has one child and is itself (a) a left child

Removing an Entry Whose Node Has One Child




FIGURE 25-7 Removing a node *N* from its parent node *P* when *N* has one child and is itself (b) a right child

Removing an Entry Whose Node Has Two Children




FIGURE 25-8 Two possible configurations of a node N that has two children

Removing an Entry Whose Node Has Two Children




FIGURE 25-9 Node N and its subtrees: (a) the entry a is immediately before the entry e , and b is immediately after e ; (b) after deleting the node that contained a and replacing e with a

Removing an Entry Whose Node Has Two Children




FIGURE 25-10 The largest entry *a* in node *N*'s left subtree occurs in the subtree's rightmost node *R*

Removing an Entry Whose Node Has Two Children





FIGURE 25-11 (a) A binary search tree;

(b) after removing *Chad*

Removing an Entry Whose Node Has Two Children

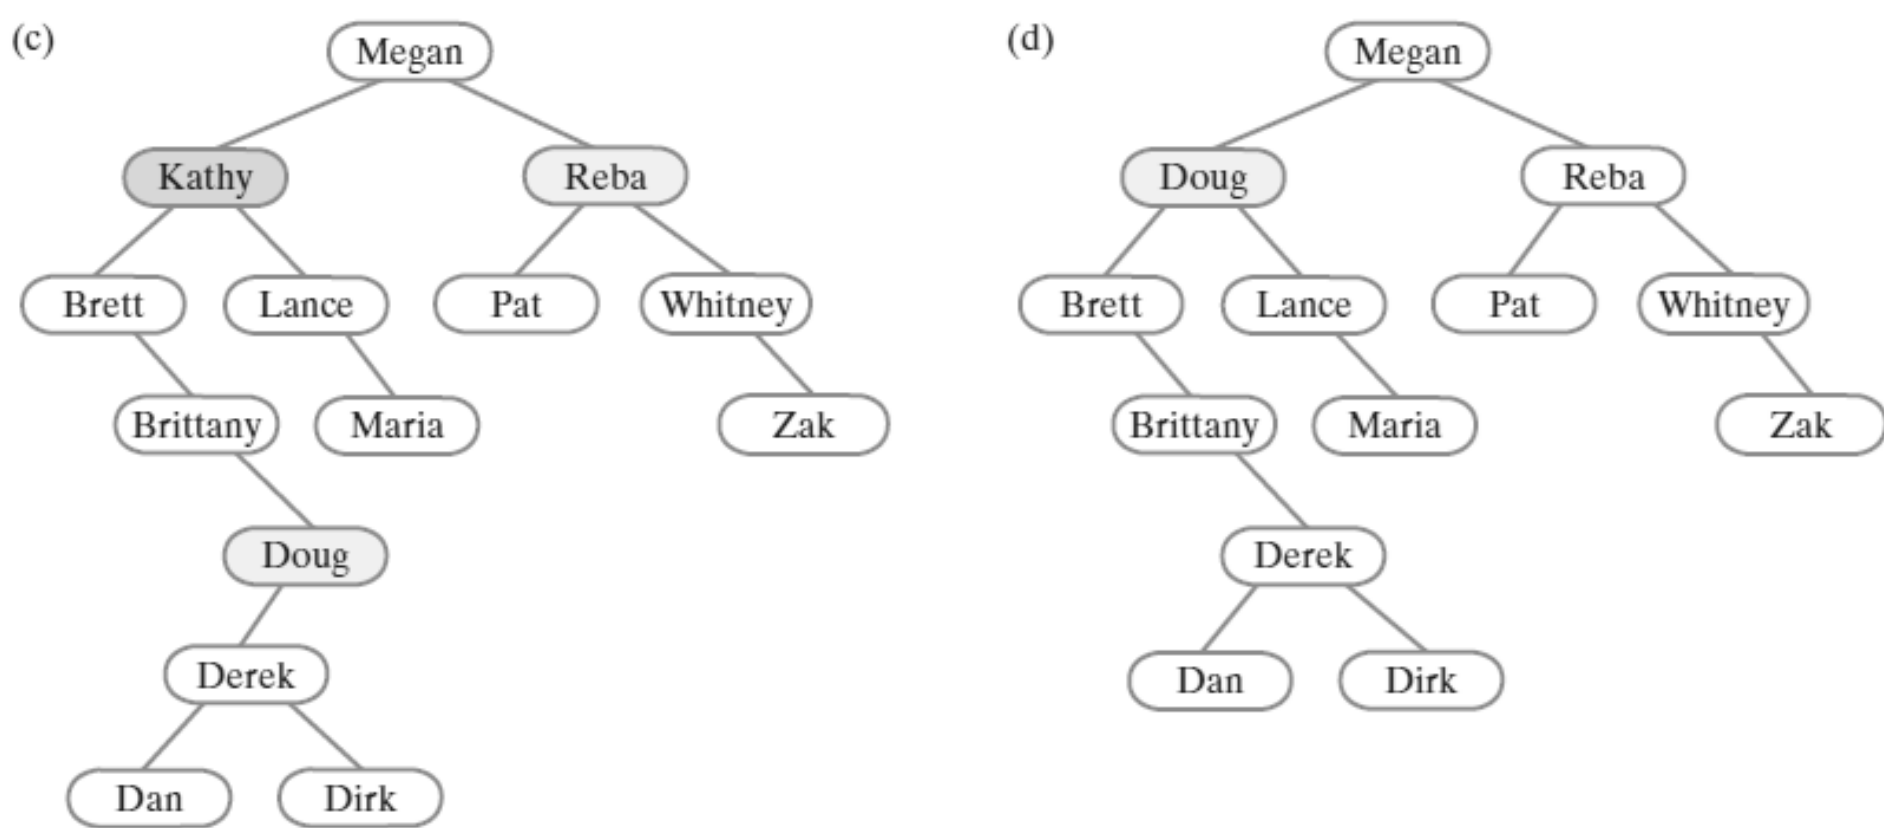


FIGURE 25-11 (c) after removing *Sean*;

(d) after removing *Kathy*

Removing an Entry in the Root




FIGURE 25-12 (a) Two possible configurations of a root that has one child; (b) after removing the root

Recursive Implementation

```
Algorithm remove(binarySearchTree, entry)
oldEntry = null
if (binarySearchTree is not empty)
{
    if (entry matches the entry in the root of binarySearchTree)
    {
        oldEntry = entry in root
        removeFromRoot(root of binarySearchTree)
    }
    else if (entry < entry in root)
        oldEntry = remove(left subtree of binarySearchTree, entry)
    else // entry > entry in root
        oldEntry = remove(right subtree of binarySearchTree, entry)
}
return oldEntry
```

Recursive algorithm describes the
method's logic at a high level

Recursive Implementation

```
// Removes an entry from the tree rooted at a given node.
// rootNode is a reference to the root of a tree.
// entry is the object to be removed.
// oldEntry is an object whose data field is null.
// Returns the root node of the resulting tree; if entry matches
//      an entry in the tree, oldEntry's data field is the entry
//      that was removed from the tree; otherwise it is null.
private BinaryNode<T> removeEntry(BinaryNode<T> rootNode, T entry,
                                   ReturnObject oldEntry)
{
    if (rootNode != null)
    {
        T rootData = rootNode.getData();
        int comparison = entry.compareTo(rootData);
        if (comparison == 0)    // entry == root entry
        {
            oldEntry.set(rootData);
            rootNode = removeFromRoot(rootNode);
        }
        else if (comparison < 0) // entry < root entry
        {
            // ... (code for left subtree) ...
        }
        else // entry > root entry
        {
            // ... (code for right subtree) ...
        }
    }
}
```

Recursive Implementation

```
else if (comparison < 0) // entry < root entry
{
    BinaryNode<T> leftChild = rootNode.getLeftChild();
    BinaryNode<T> subtreeRoot = removeEntry(leftChild, entry, oldEntry);
    rootNode.setLeftChild(subtreeRoot);
}
else // entry > root entry
{
    BinaryNode<T> rightChild = rootNode.getRightChild();
    rootNode.setRightChild(removeEntry(rightChild, entry, oldEntry));
} // end if
} // end if
return rootNode;
} // end removeEntry
```


Recursive Implementation

Algorithm removeFromRoot(rootNode)

// Removes the entry in a given root node of a subtree.

if (rootNode has two children)

{

largestNode = node with the largest entry in the left subtree of rootNode

Replace the entry in rootNode with the entry in largestNode

Remove largestNode from the tree

}

else if (rootNode has a right child)

rootNode = rootNode's right child

else

rootNode = rootNode's left child // Possibly null

// Assertion: If rootNode was a leaf, it is now null

return rootNode

Recursive Implementation

```
// Removes the entry in a given root node of a subtree.
// rootNode is the root node of the subtree.
// Returns the root node of the revised subtree.
private BinaryNode<T> removeFromRoot(BinaryNode<T> rootNode)
{
    // Case 1: rootNode has two children
    if (rootNode.hasLeftChild() && rootNode.hasRightChild())
    {
        // Find node with largest entry in left subtree
        BinaryNode<T> leftSubtreeRoot = rootNode.getLeftChild();
        BinaryNode<T> largestNode = findLargest(leftSubtreeRoot);

        // Replace entry in root
        rootNode.setData(largestNode.getData());

        // Remove node with largest entry in left subtree
        rootNode.setLeftChild(removeLargest(leftSubtreeRoot));
    }
}
```

Recursive Implementation

```
    rootNode.setData(largestNode.getData());  
    // Remove node with largest entry in left subtree  
    rootNode.setLeftChild(removeLargest(leftSubtreeRoot));  
} // end if  
  
// Case 2: rootNode has at most one child  
else if (rootNode.hasRightChild())  
    rootNode = rootNode.getRightChild();  
else  
    rootNode = rootNode.getLeftChild();  
  
// Assertion: If rootNode was a leaf, it is now null  
return rootNode;  
} // end removeEntry
```

Recursive Implementation

```
// Finds the node containing the largest entry in a given tree.  
// rootNode is the root node of the tree.  
// Returns the node containing the largest entry in the tree.  
private BinaryNode<T> findLargest(BinaryNode<T> rootNode)  
{  
    if (rootNode.hasRightChild())  
        rootNode = findLargest(rootNode.getRightChild());  
    return rootNode;  
} // end findLargest
```

Recursive Implementation

```
// Removes the node containing the largest entry in a given tree.  
// rootNode is the root node of the tree.  
// Returns the root node of the revised tree.  
private BinaryNode<T> removeLargest(BinaryNode<T> rootNode)  
{  
    if (rootNode.hasRightChild())  
    {  
        BinaryNode<T> rightChild = rootNode.getRightChild();  
        rightChild = removeLargest(rightChild);  
        rootNode.setRightChild(rightChild);  
    }  
    else  
        rootNode = rootNode.getLeftChild();  
    return rootNode;  
} // end removeLargest
```

Iterative Implementation

Algorithm remove(entry)

result = null

currentNode = *node that contains a match for entry*

parentNode = *currentNode's parent*

if (currentNode != null) // *That is, if entry is found*
{

 result = *currentNode's data (the entry to be removed from the tree)*

 // *Case 1*

 if (currentNode *has two children*)

 {

 // *Get node to remove and its parent*

 nodeToRemove = *node containing entry's inorder predecessor; it has at most one child*

 parentNode = *nodeToRemove's parent*

Copy entry from nodeToRemove to currentNode

 currentNode = nodeToRemove

 // *Assertion: currentNode is the node to be removed; it has at most one child*

 // *Assertion: Case 1 has been transformed to Case 2*

 }

 // *Case 2: currentNode has at most one child*

Delete currentNode from the tree

}

return result

Iterative Implementation

```
public T remove(T entry)
{
    T result = null;
    // Locate node (and its parent) that contains a match for entry
    NodePair pair = findNode(entry);
    BinaryNode<T> currentNode = pair.getFirst();
    BinaryNode<T> parentNode = pair.getSecond();
    if (currentNode != null) // Entry is found
    {
        result = currentNode.getData(); // Get entry to be removed
        // Case 1: currentNode has two children
        if (currentNode.hasLeftChild() && currentNode.hasRightChild())
        {
            // Replace entry in currentNode with the entry in another node
            // that has at most one child; that node can be deleted
        }
    }
}
```


Iterative Implementation

```
// Get node to remove (contains inorder predecessor; has at
// most one child) and its parent
pair = getNodeToRemove(currentNode);
BinaryNode<T> nodeToRemove = pair.getFirst();
parentNode = pair.getSecond();

// Copy entry from nodeToRemove to currentNode
currentNode.setData(nodeToRemove.getData());

currentNode = nodeToRemove;
// Assertion: currentNode is the node to be removed; it has at
// most one child
// Assertion: Case 1 has been transformed to Case 2
} // end if

// Case 2: currentNode has at most one child; delete it
removeNode(currentNode, parentNode);
} // end if

return result;
} // end remove
```


Iterative Implementation

```
private NodePair findNode(T entry)
{
    NodePair result = new NodePair();
    boolean found = false;
    . . .
    if (found)
        result = new NodePair(currentNode, parentNode);
        // Located entry is currentNode.getData()
    return result;
} // end findNode
```

The private method **findNode**

Iterative Implementation

```
// Find the inorder predecessor by searching the left subtree; it will be the largest
// entry in the subtree, occurring in the node as far right as possible
leftSubtreeRoot = left child of currentNode
rightChild = leftSubtreeRoot
priorNode = currentNode
while (rightChild has a right child )
{
    priorNode = rightChild
    rightChild = right child of rightChild
}
// Assertion: rightChild is the node to be removed and has no more than one child
```

Pseudocode for the private method **getNodeToRemove**

Iterative Implementation

```
private NodePair getNodeToRemove(BinaryNode<T> currentNode)
{
    // Find node with largest entry in left subtree by
    // moving as far right in the subtree as possible
    BinaryNode<T> leftSubtreeRoot = currentNode.getLeftChild();
    BinaryNode<T> rightChild = leftSubtreeRoot;
    BinaryNode<T> priorNode = currentNode;

    while (rightChild.hasRightChild())
    {
        priorNode = rightChild;
        rightChild = rightChild.getRightChild();
    } // end while

    // rightChild contains the inorder predecessor and is the node to
    // remove; priorNode is its parent

    return new NodePair(rightChild, priorNode);
} // end getNodeToRemove
```

Iterative Implementation

```
private void removeNode(BinaryNode<T> nodeToRemove, BinaryNode<T> parentNode)
{
    BinaryNode<T> childNode;
    if (nodeToRemove.hasLeftChild())
        childNode = nodeToRemove.getLeftChild();
    else
        childNode = nodeToRemove.getRightChild();
    // Assertion: If nodeToRemove is a leaf, childNode is null
    assert (nodeToRemove.isLeaf() && childNode == null) ||
        !nodeToRemove.isLeaf();
    if (nodeToRemove == getRootNode())
        setRootNode(childNode);
    else if (parentNode.getLeftChild() == nodeToRemove)
        parentNode.setLeftChild(childNode);
    else
        parentNode.setRightChild(childNode);
} // end removeNode
```

Efficiency of Operations

- For tree of height h
 - The operations **add**, **remove**, and **getEntry** are $O(h)$
- If tree of n nodes has height $h = n$
 - These operations are $O(n)$
- Shortest tree is full
 - Results in these operations being $O(\log n)$

Efficiency of Operations

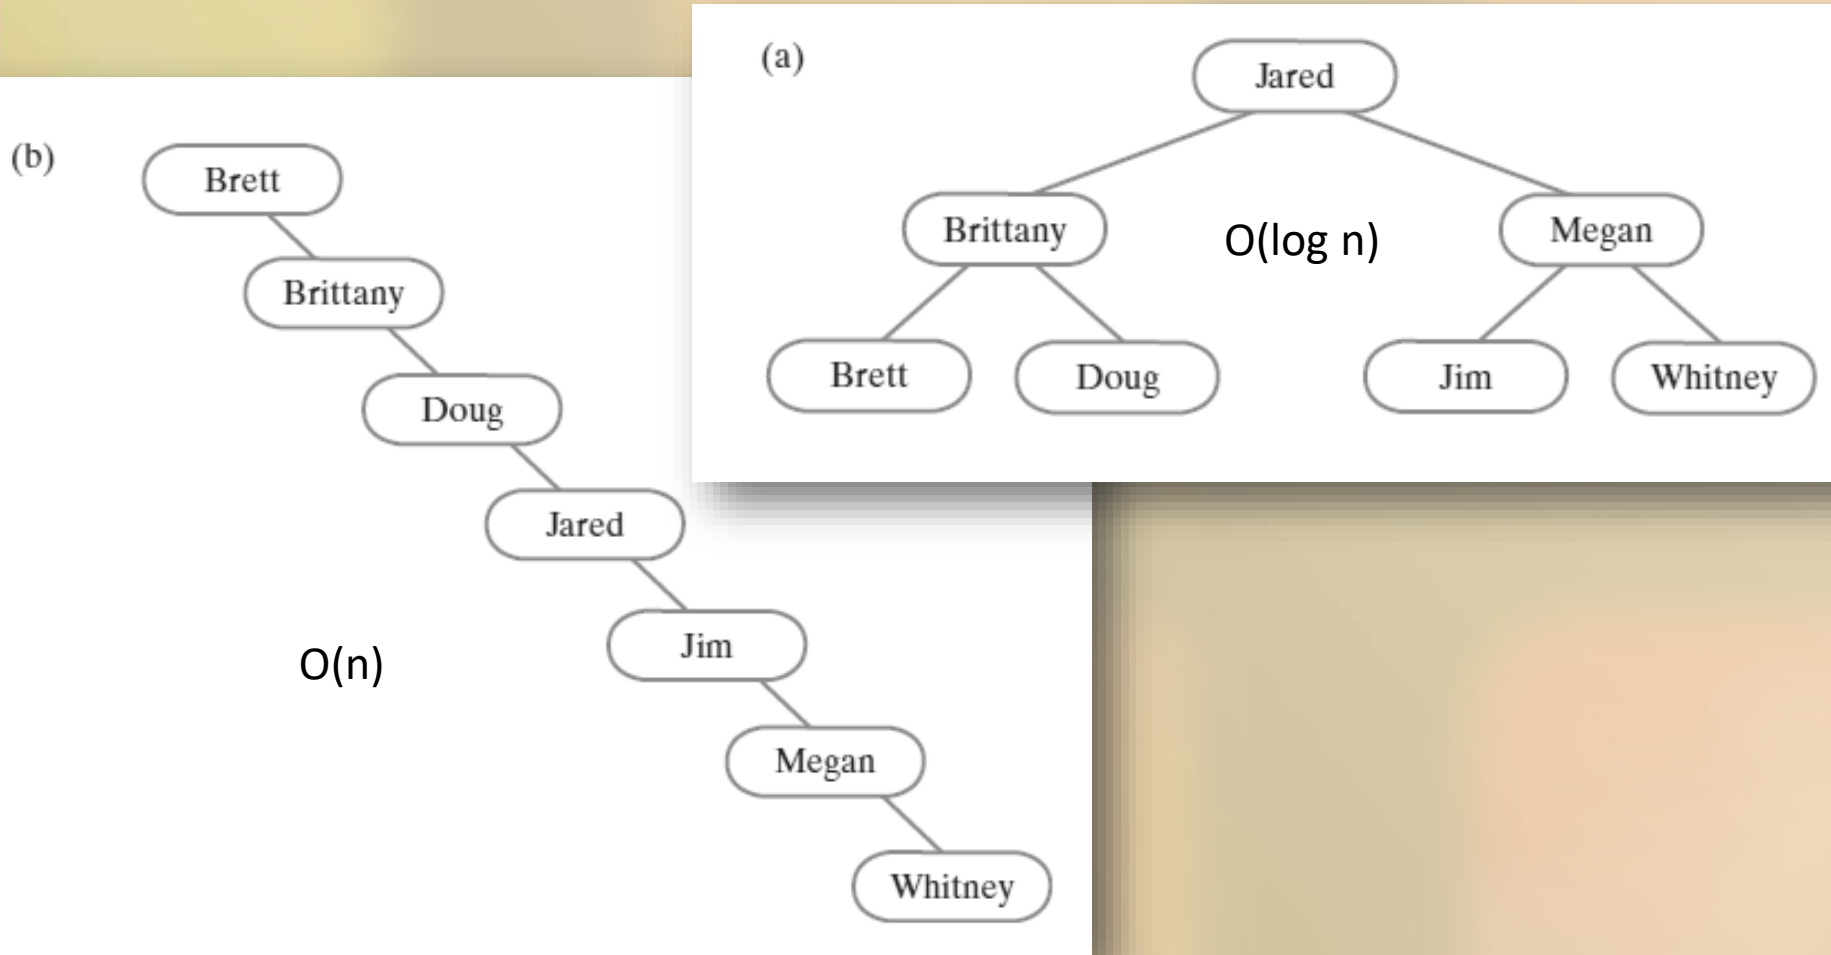


FIGURE 25-13 Two binary search trees that contain the same data

The Importance of Balance

- Do not need a full binary search tree to get $O(\log n)$ performance
- Complete tree will also give us $O(\log n)$ performance.

Order in Which Nodes Are Added

- Order in which you add entries to a binary search tree affects the shape of the tree
- If you add entries into an initially empty binary search tree, do not add them in sorted order.

End

Chapter 25